

SVMs and Feature Selection

Machine Learning (CS 545), Portland State University, Winter 2016

Bruce Marron

11 February 2016

Introduction

Support vector machines (SVMs) are supervised learning methods that prove especially useful for classification and regression tasks. The experiments described below use a linear SVM to classify email as 'spam' or 'not spam' given a feature vector in \mathbb{R}^{57} . The raw data for the experiments ("Spambase") were obtained from the University of California, Irvine (UCI) Machine Learning Repository. All data processing tasks as well as the experiments themselves were done in Python 2.7.11 (Anaconda 2.4.1 (32-bit)) using the integrated development environment (IDE), "spyder" and (primarily) the open source, "scikit-learn" machine learning package.

Data Processing

The raw data in "Spambase" contains 4601 cases of which 1813 are identified as 'spam' ('1' in Field 58) and 2788 are identified as 'not spam' ('0' in Field 58). After shuffling the subset of 'not spam' cases, 1812 of these were selected. The subset of 1812 'not spam' cases was combined with the subset of 1812 'spam' cases to create a SVM subset of "Spambase" containing an equal number of positive and negative cases. This subset was randomly shuffled and then split equally into a training subset (`tr_d`) and a test subset (`te_d`). The 57 features in all of the cases in the training subset were standardized. The 57 features in all of the cases in the test subset were likewise standardized, but using the mean and standard deviation obtained from the training subset standardization. The procedures and methods used for all data processing tasks are detailed in the script, `DataProcessing.py`.

Experiment 1

A linear SVM model was defined in Python (`sklearn.svm.SVC()`) and a 10-fold cross-validation was performed on the processed training dataset (`tr_d`) for each of 11 different C parameters ($C = 0, .1, .2, .3, .4, .5, .6, .7, .8, .9, 1$). The C parameter giving the highest accuracy ($= 0.9261$) was identified ($C = 0.6$) and the SVM model was then trained using this C parameter. The trained SVM model was applied to the test dataset (`te_d`) and accuracy, precision, recall statistics were obtained. Likewise, false positive rate (fpr) and true positive rate (tpr) statistics were determined and a receiver operating curve (ROC) was generated. Outputs from the cross-validation study are presented in Table 1, the statistics for the SVM model are presented in Table 2, and the ROC curve is presented in Figure 1. The realization of Experiment 1 is documented in the files, `Exp1_CrossValidation.py`, `Exp1_Classification.py`, and `Exp1_Classification_ROC.py`.

Table 1: The mean accuracy of a linear SVM on training data derived from the "Spambase" dataset using 10-fold cross-validation for various values of the C parameter.

Cross-Validation Run	C Parameter	Mean Accuracy
1	0.0	NA*
2	0.1	0.9211
3	0.2	0.9233
4	0.3	0.9216
5	0.4	0.9239
6	0.5	0.9244
7	0.6	0.9261
8	0.7	0.9239
9	0.8	0.9227
10	0.9	0.9216
11	1.0	0.9211

* No results obtained because of SVM error for C=0

Table 2: The accuracy, precision, and recall of a trained, linear SVM using the model parameter, $C = 0.6$. The training data were derived from the "Spambase" dataset.

SVM.ID	Accuracy	Precision	Recall
exp2.svm	0.9266	0.9239	0.9290




Figure 1: A receiver operating curve (ROC) with 310 thresholds for a trained, linear SVM using the model parameter, $C = 0.6$. The training data were derived from the "Spambase" dataset.

Experiment 2

A linear SVM model was trained on the processed training dataset (`tr_d`) using the C parameter identified as giving the highest accuracy ($C = 0.6$) from Experiment 1. After training, the weight vector for the SVM model was obtained. The absolute value of each weight was determined and the 57 features in the `tr_d` dataset were then ranked by their absolute weight from greatest value to least value. A series of SVM models was used to evaluate the test dataset (`te_d`) where, for each evaluation, the highest ranked m features ($m = 1, 2, \dots, 57$) were selected for the SVM model and all other features were set to zero. Thus, the SVM model for the first evaluation of the test dataset used only the highest ranking feature (Feature 26; $\text{abs.wgt} = 2.4654$) with the values for all other features in `te_d` set to zero. This SVM model gave an accuracy of 0.6098. The SVM model for the second evaluation of the test dataset used only the first and second highest ranking features (Feature 26; $\text{abs.wgt} = 2.4654$ and Feature 24; $\text{abs.wgt} = 1.4083$) with the values for all other features in `te_d` set to zero. This second SVM model gave an accuracy of 0.7339. The experimental results for all SVM models are given in Table 3 and are displayed graphically in Figure 2. The realization of Experiment 2 is documented in the files, `Exp2_FeatureSelection1.py` and `Exp2_FeatureSelection2.py`.




Figure 2: The cumulative effect of adding weight-ranked features on the accuracy of the SVM in Experiment 2.

Table 3: The 57 features used for the SVM in Experiment 2 ranked by the absolute value of their weight and the cumulative effect on the SVM accuracy.

Rank Order	Feature Number	SVM weight	Cumulative Accuracy
1	26	2.4654	0.6098
2	24	1.4083	0.7339
3	30	1.4016	0.7356
4	52	1.1448	0.7582
5	45	0.9517	0.7864
6	6	0.8817	0.8366
7	15	0.8394	0.8664
8	44	0.7067	0.8758
9	25	0.6905	0.8780
10	40	0.6637	0.8813
11	41	0.6196	0.8940
12	34	0.5762	0.8945

Continued on next page ...

Table 3 – continued from previous page

Rank Order	Feature Number	SVM weight	Cumulative Accuracy
13	4	0.4973	0.8984
14	23	0.4647	0.9045
15	55	0.3969	0.9161
16	22	0.3652	0.9177
17	53	0.3602	0.9210
18	38	0.3255	0.9210
19	3	0.3071	0.9227
20	28	0.2972	0.9227
21	16	0.2909	0.9199
22	7	0.2876	0.9194
23	43	0.2698	0.9199
24	54	0.2591	0.9194
25	14	0.2567	0.9194
26	32	0.2440	0.9194
27	51	0.2286	0.9183
28	56	0.2272	0.9216
29	29	0.2243	0.9232
30	20	0.2226	0.9227
31	21	0.2024	0.9221
32	27	0.1986	0.9221
33	31	0.1856	0.9210
34	18	0.1719	0.9216
35	48	0.1593	0.9238
36	8	0.1498	0.9282
37	2	0.1333	0.9177
38	10	0.1115	0.9216
39	47	0.1035	0.9216
40	17	0.1033	0.9183
41	9	0.0931	0.9210
42	19	0.0911	0.9216
43	0	0.0907	0.9249
44	39	0.0849	0.9249
45	33	0.0706	0.9249

Continued on next page ...

Table 3 – continued from previous page

Rank Order	Feature Number	SVM weight	Cumulative Accuracy
46	42	0.0680	0.9254
47	46	0.0558	0.9232
48	5	0.0471	0.9260
49	37	0.0354	0.9249
50	11	0.0301	0.9243
51	13	0.0236	0.9277
52	12	0.0234	0.9277
53	35	0.0206	0.9260
54	36	0.0187	0.9266
55	49	0.0014	0.9266
56	1	0.0003	0.9266
57	50	0.00003	0.9266

Experiment 3

Experiment 3 had exactly the same design as Experiment 2 except that features were ranked randomly rather than being ranked by the absolute value of their weight. The experimental results for all SVM models in Experiment 3 are given in Table 4 and are displayed graphically in Figure 3. The realization of Experiment 3 is documented in the files, `Exp3_FeatureSelection1.py` and `Exp3_FeatureSelection2.py`.

Table 4: The 57 features used for the SVM in Experiment 2 ranked in random order and the cumulative effect on the SVM accuracy.

Rank Order	Feature Number	SVM weight	Cumulative Accuracy
1	6	0.8817	0.6799
2	42	0.0680	0.6799
3	31	0.1856	0.6799
4	38	0.3255	0.6799
5	1	0.0003	0.6799
6	17	0.1033	0.6915
7	23	0.4647	0.7809
8	46	0.0558	0.7814

Continued on next page ...

Table 4 – continued from previous page

Rank Order	Feature Number	SVM weight	Cumulative Accuracy
9	2	0.1333	0.7775
10	27	0.1986	0.7836
11	21	0.2024	0.7880
12	44	0.7067	0.7880
13	16	0.2909	0.8035
14	48	0.1593	0.8107
15	52	1.1448	0.8399
16	28	0.2972	0.8454
17	4	0.4973	0.8509
18	0	0.0907	0.8509
19	10	0.1115	0.8548
20	18	0.1719	0.8543
21	55	0.3969	0.8592
22	19	0.0911	0.8565
23	47	0.1035	0.8620
24	51	0.2286	0.8818
25	34	0.5762	0.8796
26	33	0.0706	0.8796
27	3	0.3071	0.8796
28	53	0.3602	0.8785
29	56	0.2272	0.8807
30	40	0.6637	0.8830
31	29	0.2243	0.8835
32	11	0.0301	0.8846
33	22	0.3652	0.8912
34	49	0.0014	0.8912
35	13	0.0236	0.8912
36	24	1.4083	0.9001
37	37	0.0354	0.9023
38	50	0.00003	0.9023
39	45	0.9517	0.9028
40	26	2.4654	0.9105
41	9	0.0931	0.9100

Continued on next page ...

Table 4 – continued from previous page

Rank Order	Feature Number	SVM weight	Cumulative Accuracy
42	32	0.2440	0.9089
43	20	0.2226	0.9139
44	43	0.2698	0.9139
45	25	0.6905	0.9150
46	30	1.4016	0.9139
47	12	0.0234	0.9122
48	5	0.0471	0.9150
49	15	0.8394	0.9183
50	35	0.0206	0.9205
51	8	0.1498	0.9227
52	39	0.0849	0.9243
53	7	0.2876	0.9243
54	14	0.2567	0.9249
55	54	0.2591	0.9249
56	41	0.6196	0.9260
57	36	0.0187	0.9266

Results and Discussion

Although the accuracies for all C parameters in Experiment 1 were quite high (>0.92), there was a clear performance peak at $C=.06$. The high quality of the classifier performance is reflected in the ROC curve. Apparently, the SVM software requires a non-zero value for the C parameter.

The SVM model in Experiment 3 produced a ranked order of important features for spam detection in the "Spambase" dataset. The top five features for spam detection, given the conditions of Experiment 3, are given in Table 5. Why it is that "George" and "hp" should be included in an email is unclear, but a quick Google search confirms that these two letter sequences were, in fact, ubiquitous in the content of spam circa 2000. One might conjure various explanations for the remaining three descriptors based on legitimacy and greed. Regardless of why such descriptors are used, building SVM classifiers using features selected based on the rank of the absolute value of their weights appears to be useful. In the current case, using just the top 10 ranked features will provide nearly 90% accuracy.




Figure 3: The cumulative effect of adding randomly ranked features on the accuracy of the SVM in Experiment 3.

On the other hand, the use of randomly selected features in Experiment 4 seems, as suspected, to provide little efficacy: reaching about the same level of accuracy as in Experiment 3 would require at least 29 features. Notice also that the shape of the curve in Figure 3 lags behind that of Figure 2 in the rate at which increased accuracy is attained. Finally, both Figure 2 and Figure 3 have an asymptotic character suggesting that little gain in accuracy would be expected beyond a threshold number of features.

Table 5: The descriptors of the top five features of the "Spambase" dataset, as ranked by the absolute value of their weights.

Rank Order	Feature	Descriptor
1	26	word_freq_george
2	24	word_freq_hp
3	30	word_freq_telnet
4	52	char_freq_\$
5	45	word_freq_edu

Conclusions

Linear SVMs appear to be highly successful at the task of binary clustering. Clearly, feature selection is one of the keys to successful SVM performance. Finally, there are undoubtedly additional feature sorting algorithms which could presumably be used in conjunction with the 'ranking by absolute weight' algorithm used here to produce highly efficient classifiers.