

# Evaluating Classifiers

Reading for this topic:

T. Fawcett, An introduction to ROC analysis,  
Sections 1-4, 7

(linked from class website)

# Evaluating Classifiers

- What we want: Classifier that best predicts unseen (“test”) data
- Common assumption: Data is “iid” (independently and identically distributed)

# Topics

- Cross Validation
- Precision and Recall
- ROC Analysis
- Bias, Variance, and Noise

# Cross-Validation

# Accuracy and Error Rate

- Accuracy = fraction of correct classifications on unseen data (test set)
- Error rate =  $1 - \text{Accuracy}$

# How to use available data to best measure accuracy?

Split data into training and test sets.

But how to split?

Too little training data: Cannot learn a good model

Too little test data: Cannot evaluate learned model

Also, how to learn hyper-parameters of the model?

One solution: “ $k$ -fold cross validation”

- Used to better estimate generalization accuracy of model
- Used to learn hyper-parameters of model (“model selection”)



## Using $k$ -fold cross validation to estimate accuracy

- Each example is used both as a training instance and as a test instance.
- Instead of splitting data into “training set” and “test set”, split data into  $k$  disjoint parts:  $S_1, S_2, \dots, S_k$ .
- For  $i = 1$  to  $k$   
    Select  $S_i$  to be the “test set”. Train on the remaining data, test on  $S_i$ , to obtain accuracy  $A_i$ .
- Report  $\frac{1}{k} \sum_{i=1}^k A_i$  as the final accuracy.

# Using $k$ -fold cross validation to learn hyper-parameters (e.g., learning rate, number of hidden units, SVM kernel, etc. )

- Split data into training and test sets. Put test set aside.
- Split training data into  $k$  disjoint parts:  $S_1, S_2, \dots, S_k$ .
- Assume you are learning one hyper-parameter. Choose  $R$  possible values for this hyper-parameter.
- For  $j = 1$  to  $R$

For  $i = 1$  to  $k$

Select  $S_i$  to be the “validation set”

Train the classifier on the remaining data using the  $j$ th value of the hyperparameter

Test the classifier on  $S_i$ , to obtain accuracy  $A_{i,j}$ .

Compute the average of the accuracies:  $\bar{A}_j = \frac{1}{k} \sum_{i=1}^k A_{i,j}$

Choose the value  $j$  of the hyper-parameter with highest  $\bar{A}_j$ .

Retrain the model with all the training data, using this value of the hyper-parameter.

Test resulting model on the test set.

# Precision and Recall

# Evaluating classification algorithms

“Confusion matrix” for a given class  $c$

Actual	Predicted (or “classified”)	
	Positive (in class $c$ )	Negative (not in class $c$ )
Positive (in class $c$ )	TruePositive	FalseNegative
Negative (not in class $c$ )	FalsePositive	TrueNegative

# Example: “A” vs. “B”

Assume “A” is positive class

**Confusion Matrix**

Results from Perceptron:

<u>Instance</u>	<u>Class</u>	<u>Perception Output</u>
1	“A”	−1
2	“A”	+1
3	“A”	+1
4	“A”	−1
5	“B”	+1
6	“B”	−1
7	“B”	−1
1	“B”	−1

Actual	Predicted	
	Positive	Negative
Positive	2	2
Negative	1	3

**Accuracy:**

# Evaluating classification algorithms

“Confusion matrix” for a given class  $c$

Actual	Predicted (or “classified”)	
	Positive (in class $c$ )	Negative (not in class $c$ )
Positive (in class $c$ )	TruePositive	FalseNegative
Negative (not in class $c$ )	FalsePositive	TrueNegative


Type 2 error

Type 1 error

# Exercise 1

# Precision and Recall


**All instances in test set**






# Precision and Recall

**All instances in test set**




# Precision and Recall

**All instances in test set**



# Precision and Recall


**All instances in test set**



# Precision and Recall

$$\text{Precision} = \frac{TP}{TP + FP}$$

**All instances in test set**




# Precision and Recall

$$\text{Precision} = \frac{TP}{TP + FP}$$

**All instances in test set**

$$\text{Recall} = \frac{TP}{TP + FN}$$



Recall = Sensitivity = True Positive Rate

# Example: “A” vs. “B”

Assume “A” is positive class

Results from Perceptron:

Instance	Class	Perception Output
1	“A”	−1
2	“A”	+1
3	“A”	+1
4	“A”	−1
5	“B”	+1
6	“B”	−1
7	“B”	−1
1	“B”	−1

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

$$\text{F-measure} = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

# Creating a Precision vs. Recall Curve

$$P = \frac{TP}{TP + FP}$$

$$R = \frac{TP}{TP + FN}$$

Inst#	Class	Score	Inst#	Class	Score
1	p	.9	11	p	.4
2	p	.8	12	n	.39
3	n	.7	13	p	.38
4	p	.6	14	n	.37
5	p	.55	15	n	.36
6	p	.54	16	n	.35
7	n	.53	17	p	.34
8	n	.52	18	n	.33
9	p	.51	19	p	.30
10	n	.505	20	n	.1

**Results of classifier**

Threshold	Accuracy	Precision	Recall
.9	11/20	1	1/10
.8	12/20	1	2/10
.7			
.6			
.5			
.4			
.3			
.2			
.1	10/20	10/20	1
$-\infty$			

# ROC Analysis



# Receiver Operating Characteristic (ROC) Curves


- Alternative to precision/recall curves
- Shows tradeoff between true positive rate and false positive rate.

True positive rate =  $TP / (TP + FN)$

(“sensitivity”)

False positive rate =  $FP / (TN + FP)$

(1 - “specificity”)



# Creating a ROC Curve

$$\text{True Positive Rate (= Recall)} = \frac{TP}{TP + FN}$$

$$\text{False Positive Rate} = \frac{FP}{TN + FP}$$

Inst#	Class	Score	Inst#	Class	Score
1	p	.9	11	p	.4
2	p	.8	12	n	.39
3	n	.7	13	p	.38
4	p	.6	14	n	.37
5	p	.55	15	n	.36
6	p	.54	16	n	.35
7	n	.53	17	p	.34
8	n	.52	18	n	.33
9	p	.51	19	p	.30
10	n	.505	20	n	.1

**Results of classifier**

Threshold	Accuracy	TPR	FPR
.9		1/10	0
.8		2/10	0
.7		2/10	1/10
.6			
.5			
.4			
.3			
.2			
.1		1	1
$-\infty$			

## Area under ROC curve (AUC)

- Summary statistic: Area under ROC curve (AUC) = probability that classifier will rank a randomly chosen positive instance higher than a randomly chosen negative instance.
- AUC is always between 0 and 1.

# How to create a ROC curve for a perceptron

- Run your classifier on each instance in the test data, without doing the *sgn* step:

$$Score(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x}$$

- Get the range  $[min, max]$  of your scores
- Divide the range into about 200 thresholds, including  $-\infty$  and  $max$
- For each threshold, calculate TPR and FPR
- Plot TPR (y-axis) vs. FPR (x-axis)


# In-Class Exercise 1

# Bias, Variance, and Noise

# Bias:

Classifier is not powerful enough to represent the true function;  
that is, it *underfits* the function


.






# Variance:

Classifier's hypothesis depends on specific training set; that is, it *overfits* the function



## Noise:

Underlying process generating data is stochastic, or data has errors or outliers



- Examples of bias?
- Examples of variance?
- Examples of noise?



# Bias/variance tradeoff

- Models with **too many parameters** may fit the training data well (**low bias**), but are sensitive to choice of training set (**high variance**)


- Generalization error is due to **overfitting**

- Models with **too few parameters** may not fit the data well (**high bias**) but are consistent across different training sets (**low variance**)

- Generalization error is due to **underfitting**



# Illustration of Bias / Variance Tradeoff



Hastie, Tibshirani, Friedman "Elements of Statistical Learning" 2001

# Model Error Decomposition: The Math

Let  $h(x)$  be our learned model, which estimates true function  $f(x)$ .

$$Error(x) = E\left[\left(f(x) - h(x)\right)^2\right]$$

$$= \left(E[h(x)] - f(x)\right)^2 + E\left[h(x) - E[h(x)]\right]^2 + \sigma_e^2$$

$$= \text{bias}^2 + \text{variance}^2 + \text{irreducible error}$$

## In-Class Exercise 2