

SVMs, Part 2

- Summary of SVM algorithm
- Examples of “custom” kernels
- Standardizing data for SVMs
- Soft-margin SVMs

Summary of SVM algorithm

Given training set

$$S = \{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_m, t_m) \mid (\mathbf{x}_k, t_k) \in \mathfrak{R}^n \times \{+1, -1\}\}$$

1. Choose a kernel function $K(\mathbf{x}, \mathbf{z})$.
2. Apply optimization procedure (using the kernel function K) to find support vectors \mathbf{x}_k , coefficients α_k , and bias b .
3. Given a new instance, \mathbf{x} , find the classification of \mathbf{x} by computing

$$\text{class}(\mathbf{x}) = \text{sgn} \left(\left(\sum_{k \in \text{support vectors}} \alpha_k k(\mathbf{x}, \mathbf{x}_k) \right) + b \right)$$

Most commonly used kernels

- Linear

$$K(\mathbf{x}, \mathbf{x}_i) = \mathbf{x} \cdot \mathbf{x}_i$$

- Polynomial

$$K(\mathbf{x}, \mathbf{x}_i) = [(\mathbf{x} \cdot \mathbf{x}_i) + 1]^d$$

- Gaussian (or “radial basis function”)

$$K(\mathbf{x}, \mathbf{x}_i) = e^{-\gamma|\mathbf{x}-\mathbf{x}_i|^2}$$

- Sigmoid

$$K(\mathbf{x}, \mathbf{x}_i) = \tanh(a\mathbf{x} \cdot \mathbf{x}_i + b)$$

How to define your own kernel

- Given training data ($\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m$)
- Algorithm for SVM learning uses *kernel matrix* (also called *Gram matrix*):

$$\mathbf{K}_{i,j} = K(\mathbf{x}_i, \mathbf{x}_j), \text{ for } i, j = 1, \dots, m$$

$\mathbf{K} =$

$K(\mathbf{x}^1, \mathbf{x}^1)$	$K(\mathbf{x}^1, \mathbf{x}^2)$	$K(\mathbf{x}^1, \mathbf{x}^3)$...	$K(\mathbf{x}^1, \mathbf{x}^m)$
$K(\mathbf{x}^2, \mathbf{x}^1)$	$K(\mathbf{x}^2, \mathbf{x}^2)$	$K(\mathbf{x}^2, \mathbf{x}^3)$		$K(\mathbf{x}^2, \mathbf{x}^m)$
...
$K(\mathbf{x}^m, \mathbf{x}^1)$	$K(\mathbf{x}^m, \mathbf{x}^2)$	$K(\mathbf{x}^m, \mathbf{x}^3)$...	$K(\mathbf{x}^m, \mathbf{x}^m)$

How to define your own kernel

- We can choose some function K , and compute the kernel matrix \mathbf{K} using the training data.
- We just have to guarantee that our kernel defines an inner product on some feature space
- Mercer's Theorem: : If K is “symmetric positive semidefinite”, it defines a kernel, that is, it defines an inner product in some feature space.
- We don't even have to know what that feature space is!
- \mathbf{K} is symmetric if $\mathbf{K} = \mathbf{K}^T$
- \mathbf{K} is semidefinite if all the eigenvalues of \mathbf{K} are non-negative.

From www.cs.pitt.edu/~tomas/cs3750/kernels.ppt:

- Design criteria - we want kernels to be
 - **valid** – Satisfy Mercer condition of positive semidefiniteness
 - **good** – embody the “true similarity” between objects
 - **appropriate** – generalize well
 - **efficient** – the computation of $K(\mathbf{x}, \mathbf{x}')$ is feasible

Example of Simple “Custom” Kernel

Similarity between DNA Sequences:

E.g.,

$s_1 = \text{GAATGTCCTTTCTCTAAGTCCTAAG}$

$s_2 = \text{GGAGACTTACAGGAAAGAGATTTCG}$

Define “Hamming Distance Kernel”:

$\text{hamming}(s_1, s_2) = \text{number of sites where strings match}$

Kernel matrix for *hamming* kernel

Suppose training set is

$s_1 = \text{GAATGTCCTTTCTCTAAGTCCTAAG}$

$s_2 = \text{GGAGACTTACAGGAAAGAGATTTCG}$

$s_3 = \text{GGAAACTTTCGGGAGAGAGTTTCG}$

What is the Kernel matrix **K**?

K	s_1	s_2	s_3
s_1			
s_2			
s_3			

In-class exercises

Data Standardization

As we do for neural networks, we need to do data standardization for SVMs to avoid imbalance among feature scales:

Let μ_i denote the mean value of feature i in the training data, and σ_i denote the corresponding standard deviation. For each training example \mathbf{x} , replace each x_i as follows:

$$x'_i = \frac{x_i - \mu_i}{\sigma_i}$$

Scale the test data in the same way, using the μ_i and σ_i values computed from the training data, **not** the test data.

Hard- vs. soft- margin SVMs

Hard-margin SVMs


Find \mathbf{w} and b by doing the following minimization:

$$\min_{\mathbf{w}, b} \left(\frac{1}{2} \|\mathbf{w}\|^2 \right)$$

subject to:


$$t_k (\mathbf{w} \cdot \mathbf{x}_k + b) \geq 1, \quad k = 1, \dots, m$$

$$(t_k \in \{-1, +1\})$$



Extend to **soft-margin** SVMs

Allow some instances to be misclassified, or fall within margins, but penalize them by distance to margin hyperplane



<http://nlp.stanford.edu/IR-book/html/htmledition/img1260.png>

Revised optimization problem:

Find \mathbf{w} and b by doing the following minimization:

$$\min_{\mathbf{w}, b} \left(\frac{1}{2} \|\mathbf{w}\|^2 \right) + C \sum_k \xi_k$$

subject to:

$$t_k (\mathbf{w} \cdot \mathbf{x}_k + b) \geq 1 - \xi_k, \quad k = 1, \dots, m$$

$$(t_k \in \{-1, +1\})$$

Optimization tries to keep ξ_k 's to zero while maximizing margin.

C is parameter that trades off margin width with misclassifications

Why use soft-margin SVMs?

- Always can be optimized (unlike hard-margin SVMs)
- More robust to outliers, noise
- **However:** Have to set C parameter

SVM parameters to set

- Kernel function (and associated parameters)
- C (tradeoff between margin length and misclassifications)

Demo with SVM_light

Quiz 2

Thursday Jan. 21

Time allotted: 30 minutes.

Format: You are allowed to bring in one (double-sided) page of notes to use during the quiz. You may bring/use a calculator, but you don't really need one.

What you need to know for quiz

Multilayer Neural Networks:

- Definition of sigmoid activation function
- How forward propagation works
- Definition of *momentum* in weight updates

Support vector machines

- Definition of margin
- Definition of “support vector”
- Given support vectors, alphas, and bias, be able to compute weight vector and find equation of separating hyperplane
- Describe what a kernel function is and what its role is in SVMs (to the level described in class).
- Describe what a Kernel Matrix is, and be able to compute one, given a training set and a kernel function.
- Describe what the inputs to the SVM algorithm are, and what its outputs.
- Understand solutions to in-class exercises